

CSCI 145 Problem Set 9

October 28, 2025

Submission Instructions

Please upload *your* work by **11:59pm Monday November 3, 2025**.

- You are encouraged to discuss ideas and work with your classmates. However, you **must acknowledge** your collaborators at the top of each solution on which you collaborated with others and you **must write** your solutions independently.
- Your solutions to theory questions must be written legibly, or typeset in LaTeX or markdown. If you would like to use LaTeX, you can import the source of this document (available from the course webpage) to Overleaf.
- I recommend that you write your solutions to coding questions in a Jupyter notebook using Google Colab.
- You should submit your solutions as a **single PDF** via the assignment on Gradescope.

Grading: The point of the problem set is for *you* to learn. To this end, I hope to disincentivize the use of LLMs by **not** grading your work for correctness. Instead, you will grade your own work by comparing it to my solutions. This self-grade is due the Friday *after* the problem set is due, also on Gradescope.

Problem 1: Autoencoders and Variational Autoencoders

In this problem, we will explore unsupervised learning through autoencoders (AEs) and variational autoencoders (VAEs).

Part A: KL Divergence from First Principles

Let P be the univariate normal distribution $\mathcal{N}(\mu, \sigma^2)$ and Q be the univariate normal distribution $\mathcal{N}(0, 1)$. Starting from the scalar densities

$$p(z) = \frac{1}{\sqrt{2\pi}\sigma} \exp\left(-\frac{(z-\mu)^2}{2\sigma^2}\right), \quad q(z) = \frac{1}{\sqrt{2\pi}} \exp\left(-\frac{z^2}{2}\right),$$

derive the KL-divergence

$$\mathbb{E}_{z \sim P} \left[\frac{\log p(z)}{\log q(z)} \right] = \frac{1}{2} (\mu^2 + \sigma^2 - 1 - \log \sigma^2).$$

Hint: Take logs, and notice that $\mathbb{E}[(z - \mu)^2] = \sigma^2$ and $\mathbb{E}[z^2] = \sigma^2 + \mu^2$.

What is the KL divergence when we have a *multivariate* distribution where each dimension is independent? That is,

$$p(\mathbf{z}) = \prod_{i=1}^d \frac{1}{\sqrt{2\pi}\sigma_i} \exp\left(-\frac{(z_i - \mu_i)^2}{2\sigma_i^2}\right), \quad q(\mathbf{z}) = \prod_{i=1}^d \frac{1}{\sqrt{2\pi}} \exp\left(-\frac{z_i^2}{2}\right).$$

Part B: Autoencoder

Implement and train an autoencoder with a bottleneck $k = 2$ on MNIST.

Once trained, sample a uniform grid of 25 points from the latent dimension in the domain $[0, 1] \times [0, 1]$. Decode those points with the decoder and plot the results. How do they look?

Part C: Variational Autoencoder

Implement a VAE, also with bottleneck $k = 2$. In particular, you should produce $\mu_{\mathbf{x}} \in \mathbb{R}^2$ and $\Sigma^2 = \begin{bmatrix} \sigma_1^2 & 0 \\ 0 & \sigma_2^2 \end{bmatrix}$. Train the VAE with the reconstruction and multivariate KL loss you derived in part A.

Like for the autoencoder, plot the 25 decoded points. How do those decoded with the VAE compare to those decoded with the autoencoder?